



DEPARTMENT OF COMPUTER SCIENCE

CircuitFlow: A Domain Specific Language for Dataflow Programming

Riley Evans

A dissertation submitted to the University of Bristol in accordance with the requirements of the degree
of Master of Engineering in the Faculty of Engineering.

Thursday 6th May, 2021

Declaration

This dissertation is submitted to the University of Bristol in accordance with the requirements of the degree of MEng in the Faculty of Engineering. It has not been submitted for any other degree or diploma of any examining body. Except where specifically acknowledged, it is all the work of the Author.

Riley Evans, Thursday 6th May, 2021

Contents

1	Introduction	1
2	Background	3
2.1	Dataflow Programming	3
2.1.1	The Benefits	4
2.1.2	Dataflow Diagrams	4
2.1.3	Kahn Process Networks (KPNs)	5
2.2	Domain Specific Languages (DSLs)	5
2.2.1	Deep Embeddings	5
2.2.2	Shallow Embeddings	6
2.3	Higher Order Functors	6
2.4	Data types à la carte	8
2.5	Dependently Typed Programming	9
2.5.1	DataKinds Language Extension	9
2.5.2	Singletons	10
2.5.3	Type Families	10
2.5.4	Summary	10
2.6	Heterogeneous Lists	10
2.7	Existential Types	10
2.8	Phantom Type Parameters	11
2.9	Monadic Resource Theories	11
3	The Language	13
3.1	Language Requirements	13
3.2	Tasks	13
3.2.1	Data Stores	13
3.2.2	Task Constructor	14
3.3	First Attempt: Chains	14
3.3.1	Trees as Chains	15
3.3.2	Evaluation	16
3.4	Final Attempt: Circuit	17
3.4.1	Constructors	17
3.4.2	Combined Data Stores	19
3.4.3	Multi-Input Tasks	20
3.4.4	mapC operator	21
3.4.5	Completeness	21
3.4.6	Evaluation	21
4	Implementation	23
4.1	Requirements	23
4.2	Circuit AST	23
4.2.1	IFunctor	23
4.2.2	Type-Indexed Data Types à la Carte	24
4.3	Process Network	25
4.3.1	Network Typeclass	25
4.3.2	The Basic Network Representation	25
4.4	Translation to a Network	26

4.4.1	Indexed Monadic Catamorphism — icataM	26
4.4.2	BuildNetworkAlg	27
4.4.3	The Translation	27
4.5	UIDS	28
4.6	Failure in the Process Network	28
4.6.1	Maybe Monad	28
4.6.2	Except Monad	28
5	Examples	29
5.1	Song Data Aggregation	29
5.1.1	Building the pre-processing Circuit	29
5.1.2	Building the prediction Circuit	31
5.2	lhs2TeX Build System	32
5.2.1	Building the Circuit	32
5.2.2	Using the Circuit	33
5.3	Types saving the day	34
6	Critical Evaluation	37
6.1	Runtime comparison	37
6.1.1	Lazy evaluation problems	37
6.2	Use of Library	37
6.2.1	Other Libraries	37
6.3	Type safety	37
7	Conclusion	39

List of Figures

2.1	Luigi dependency graph [10]	3
2.2	Quartz composer [4]	4
2.3	An example dataflow and its imperative approach.	4
2.4	A sequence of node firings in a KPN	5
3.1	A Pipe (a) and its corresponding dataflow diagram (b).	16
3.2	The constructors in the Circuit library alongside their graphical representation.	17
3.3	A graphical representation of a task with multiple dependencies	20
5.1	A dataflow diagram for pre-processing the song data	29
5.2	A Broken Luigi Example	35
5.3	A Broken Circuit Example	35

Notation and Acronyms

DSL Domain Specific Language

EDSL Embedded DSL

FIFO First-In First-Out

KPN Kahn Process Network

GPL General Purpose Language

DPN Data Process Network

DAG Directed Acyclic Graph

AST Abstract Syntax Tree

PID Process Identifier

Acknowledgements

Change this to something meaningful

It is common practice (although totally optional) to acknowledge any third-party advice, contribution or influence you have found useful during your work. Examples include support from friends or family, the input of your Supervisor and/or Advisor, external organisations or persons who have supplied resources of some kind (e.g., funding, advice or time), and so on.

Chapter 1

Introduction

Write an introduction (do near the end)

Chapter 2

Background

2.1 Dataflow Programming

Dataflow programming is a paradigm that models applications as a directed graph. The nodes of the graph have inputs and outputs and are pure functions, therefore have no side effects. It is possible for a node to be a: source; sink; or processing node. A source is a read-only storage: it can be used to feed inputs into processes. A sink is a write-only storage: it can be used to store the outputs of processes. Processes will read from either a source or the output of another process, and then produce a result which is either passed to another process or saved in a sink. Edges connect these nodes together, and define the flow of information.

Example - Data Pipelines A common use of dataflow programming is in pipelines that process data. This paradigm is particularly helpful as it helps the developer to focus on each specific transformation on the data as a single component. Avoiding the need for long and laborious scripts that could be hard to maintain. One example of a data pipeline tool that makes use of dataflow programming is Luigi [19]. An example dataflow graph produced by the tool is shown in Figure 2.1




Figure 2.1: Luigi dependency graph [10]

Example - Quartz Composer Apple developed a tool included in XCode, named Quartz Composer, which is a node-based visual programming language [1]. As seen in Figure 2.2, it uses a visual approach to programming connecting nodes with edges. This allows for quick development of programs that process and render graphical data, without the user having to write a single line of code. This means that even non-programmers are able to use the tool.

Example - Spreadsheets A widely used example of dataflow programming is in spreadsheets. A cell in a spreadsheet can be thought of as a single node. It is possible to specify dependencies to other cells through the use of formulas. Whenever a cell is updated it sends its new value to those who depend on it, and so on. Work has also been done to visualise spreadsheets using dataflow diagrams, to help debug ones that are complex [8].



Figure 2.2: Quartz composer [4]

2.1.1 The Benefits

Visual The dataflow paradigm uses graphs, which make programming visual. It allows the end-user programmer to see how data passes through the program, much easier than in an imperative approach. In many cases, dataflow programming languages use drag and drop blocks with a graphical user interface to build programs. For example, Tableau Prep [23], that makes programming more accessible to users who do not have programming skills.

Implicit Parallelism Moore's law states that the number of transistors on a computer chip doubles every two years [16]. This meant that the chips' processing speeds also increased in alignment with Moore's law. However, in recent years this is becoming harder for chip manufacturers to achieve [2]. Therefore, chip manufactures have had to turn to other approaches to increase the speed of new chips, such as multiple cores. It is this approach the dataflow programming can effectively make use of. Since each node in a dataflow is a pure function, it is possible to parallelise implicitly. No node can interact with another node, therefore there are no data dependencies outside of those encoded in the dataflow. Thus eliminating the ability for a deadlock to occur.

2.1.2 Dataflow Diagrams

Dataflow programs are typically viewed as a graph. An example dataflow graph along with its corresponding imperative approach, can be found in Figure 2.3. The nodes 100, X , and Y are sources as they are only read from. C is a sink as it is wrote to. The remaining nodes are all processes, as they have some number of inputs and compute a result.




Figure 2.3: An example dataflow and its imperative approach.

In this diagram is possible to see how implicit parallelisation is possible. Both A and B can be calculated simultaneously, with C able to be evaluated after they are complete.

2.1.3 Kahn Process Networks (KPNs)

A method introduced by Gilles Kahn, Kahn Process Networks (KPNs) realised the concept of dataflow networks through the use of threads and unbounded First-In First-Out (FIFO) queues [11]. The FIFO queue is one where the items are output in the same order that they are added. A node in the dataflow becomes a thread in the process network. Each FIFO queue represents the edges connecting the nodes in a graph. The threads are then able to communicate through FIFO queues. The node can have multiple input queues and is able to read any number of values from them. It will then compute a result and add it to an output queue. Kahn imposed a restriction on a process in a KPNs that the thread is suspended if it attempts to fetch a value from an empty queue. The thread is not allowed to test for the presence of data in a queue.




Figure 2.4: A sequence of node firings in a KPN

Parks described a variant of KPNs, called Data Process Networks (DPNs) [14]. They recognise that if functions have no side effects then they have no values to be shared between each firing. Therefore, a pool of threads can be used with a central scheduler instead.

2.2 Domain Specific Languages (DSLs)

A DSL is a programming language that has a specialised domain or use-case. This differs from a General Purpose Language (GPL), which can be applied across a larger set of domains, and are generally turing complete. HTML is an example of a DSL: it is good for describing the appearance of websites, however, it cannot be used for more generic purposes, such as adding two numbers together.

Approaches to Implementation DSLs are typically split into two categories: standalone and embedded. Standalone DSLs require their own compiler and typically their own syntax; HTML would be an example of a standalone DSL. Embedded DSLs (EDSLs) use an existing language as a host, therefore they use the syntax and compiler from the host. This means that they are easier to maintain and often quicker to develop than standalone DSLs. An EDSL, can be implemented using two differing techniques: deep and shallow embeddings.

2.2.1 Deep Embeddings

A deep embedding is when the terms of the DSL will construct an Abstract Syntax Tree (AST) as a host language datatype. Semantics can then be provided later on with evaluation functions. Consider the example of a minimal non-deterministic parser combinator library [26].

```
data Parserd (a :: Type) where
  Satisfyd :: (Char → Bool) → Parserd Char
  Ord      :: Parserd a → Parserd a → Parserd a
```

This can be used to build a parser that can parse the characters 'a' or 'b'.

```

aorbd :: Parserd Char
aorbd = Satisfyd (≡ 'a') `Ord` Satisfyd (≡ 'b')

```

However, this parser does not have any semantics, therefore this needs to be provided by the evaluation function `parse`.

```

parsed :: Parserd a → String → [(a, String)]
parsed (Satisfyd p) = λcase
  []      → []
  (t : ts') → [(t, ts') | p t]
parsed (Ord px py) = λts → parsed px ts ++ parsed py ts

```

The program can then be evaluated by the `parsed` function. For example, `parsed aorbd "a"` evaluates to `[('a', "")]`, and `parsed aorbd "c"` evaluates to `[]`.

A key benefit for deep embeddings is that the structure can be inspected, and then modified to optimise the user code: Parsley [25] makes use of such techniques to create optimised parsers. Another benefit, is that you can provide multiple interpretations, by specifying different evaluation functions. However, they also have drawbacks - it can be laborious to add a new constructor to the language. Since it requires that all functions that use the deep embedding be modified to add a case for the new constructor [21].

2.2.2 Shallow Embeddings

In contrast, a shallow approach is when the terms of the DSL are defined as first class components of the language. For example, a function in Haskell. Components can then be composed together and evaluated to provide the semantics of the language. Again a simple parser example can be considered.

```

newtype Parsers a = Parsers { parses :: String → [(a, String)] }
ors :: Parsers a → Parsers a → Parsers a
ors (Parsers px) (Parsers py) = Parsers (λts → px ts ++ py ts)
satisfys :: (Char → Bool) → Parsers Char
satisfys p = Parsers (λcase
  []      → []
  (t : ts') → [(t, ts') | p t])

```

The same `aorbs` parser can be constructed from these functions, avoiding the need for an intermediate AST.

```

aorbs :: Parsers Char
aorbs = satisfys (≡ 'a') `ors` satisfys (≡ 'b')

```

Using a shallow implementation has the benefit of being able add new ‘constructors’ to a DSL, without having to modify any other functions. Since each ‘constructor’, produces the desired result directly. However, this causes one of the main disadvantages of a shallow embedding - the structure cannot be inspected. This means that optimisations cannot be made to the structure before evaluating it.

2.3 Higher Order Functors

It is possible to capture the shape of an abstract datatype as a `Functor`. The use of a `Functor` allows for the specification of where a datatype recurses. Consider an example on a small expression language:

```

data Expr = Add Expr Expr
          | Val Int

```

The recursion within the `Expr` datatype can be removed to form `ExprF`. The recursive steps can then be specified in the `Functor` instance.

```

data ExprF f = AddF f f
              | ValF Int

```

```
instance Functor ExprF where
  fmap f (AddF x y) = AddF (f x) (f y)
  fmap f (ValF x)   = ValF x
```

To regain a datatype that is isomorphic to the original datatype, the recursive knot need to be tied. This can be done with `Fix`, to get the fixed point of `ExprF`:

```
data Fix f = In (f (Fix f))
type Expr' = Fix ExprF
```

There is, however, one problem: a `Functor` expressing the a parser language is required to be typed. Parsers require the type of the tokens being parsed. For example, a parser reading tokens that make up an expression could have the type `Parser Expr`. A `Functor` does not retain this type information needed in a parser.

IFunctors Instead a type class called `IFunctor` — also known as `HFunctor` — can be used, which is able to maintain the type indicies [15]. This makes use of \sim , which represents a natural transformation from f to g . `IFunctor` can be thought of as a functor transformer: it is able to change the structure of a functor, whilst preserving the values inside it [13]. Whereas a functor changes the values inside a structure.

```
type (~) f g =  $\forall a. f a \rightarrow g a$ 
class IFunctor iF where
  imap :: (f  $\sim$  g)  $\rightarrow$  iF f  $\sim$  iF g
```

The shape of `Parser` can be seen in `ParserF` where the f marks the recursive spots. The type f represents the type of the children of that node. In most cases this will be itself.

```
data ParserF (f :: *  $\rightarrow$  *) (a :: *) where
  SatisfyF :: (Char  $\rightarrow$  Bool)  $\rightarrow$  ParserF f Char
  OrF      :: f a  $\rightarrow$  f a  $\rightarrow$  ParserF f a
```

An `IFunctor` instance can be defined, which follow the same structure as a standard `Functor` instance.

```
instance IFunctor ParserF where
  imap _ (SatisfyF s) = SatisfyF s
  imap f (OrF px py) = OrF (f px) (f py)
```

`Fix` is used to get the fixed point of a `Functor`, to get the indexed fixed point `IFix` can be used.

```
newtype IFix iF a = In (iF (IFix iF) a)
```

The fixed point of `ParserF` is `Parser3`.

```
type Parserfixed = IFix ParserF
```

In a deep embedding, the **AST** can be traversed and modified to make optimisations, however, it may not be the best representation when evaluating it. This means that it might be transformed to a different representation. In the case of a parser, this could be a stack machine. Now that the recursion in the datatype has been generalised, it is possible to create a mechanism to perform this transformation. An indexed *catamorphism* is one such way to do this, it is a generalised way of folding an abstract datatype. The use of a catamorphism removes the recursion from any folding of the datatype. This means that the algebra can focus on one layer at a time. This also ensures that there is no re-computation of recursive calls, as this is all handled by the catamorphism. The commutative diagram below describes how to define a catamorphism, that folds an `IFix iF a` to a `f a`.

$$\begin{array}{ccc}
\text{iF (IFix iF) a} & \xrightarrow{\text{imap (icata alg)}} & \text{iF f a} \\
\text{inop} \uparrow \downarrow \text{In} & & \downarrow \text{alg} \\
\text{IFix iF a} & \xrightarrow{\text{icata alg}} & \text{f a}
\end{array}$$

`icata` is able to fold an `IFix iF a` and produce an item of type `f a`. It uses the algebra argument as a specification of how to transform a single layer of the datatype.

```
icata :: IFunctor iF => (iF f ~> f) -> IFix iF ~> f
icata alg (lIn x) = alg (imap (icata alg) x)
```

The resulting type of `icata` is `f a`, therefore the `f` has kind `* -> *`. This could be `IFix ParserF`, which would be a transformation to the same structure, possibly applying optimisations to the [AST](#).

2.4 Data types à la carte

When building a [DSL](#) one problem that becomes quickly prevalent, the so called *Expression Problem* [\[24\]](#). The expression problem is a trade off between a deep and shallow embedding. In a deep embedding, it is easy to add multiple interpretations to the [DSL](#) - just add a new evaluation function. However, it is not easy to add a new constructor, since all functions will need to be modified to add a new case for the constructor. The opposite is true in a shallow embedding.

One possible attempt at fixing the expression problem is data types à la carte [\[22\]](#). It combines constructors using the co-product of their signatures. This is defined as:

```
data (f :+: g) a = L (f a) | R (g a)
```

It is also the case that if both `f` and `g` are `Functors` then so is `f :+: g`.

```
instance (Functor f, Functor g) => Functor (f :+: g) where
  fmap f (L x) = L (fmap f x)
  fmap f (R y) = R (fmap f y)
```

For each constructor it is possible to define a new data type and a `Functor` instance specifying where it recurses.

```
data ValF2 f = ValF2 Int
data MulF2 f = MulF2 f f
instance Functor ValF2 where
  fmap f (ValF2 x) = ValF2 x
instance Functor MulF2 where
  fmap f (MulF2 x y) = MulF2 (f x) (f y)
```

By using `Fix` to tie the recursive knot, the `Fix (Val :+: Mul)` data type would be isomorphic to the original `Expr` datatype found in [Section 2.3](#).

One problem that now exist, however, is that it is now rather difficult to create expressions, take a simple example of 12×34 .

```
exampleExpr :: Fix (ValF2 :+: MulF2)
exampleExpr = In (R (MulF2 (In (L (ValF2 12))) (In (L (ValF2 34)))))
```

It would be beneficial if there was a way to add these `Ls` and `Rs` automatically. Fortunately there is a method using injections. The `:<` type class captures the notion of subtypes between `Functors`.

```
class (Functor f, Functor g) => f :<: g where
  inj :: f a -> g a
instance Functor f => f :<: f where
  inj = id
instance (Functor f, Functor g) => f :<: (f :+: g) where
  inj = L
instance (Functor f, Functor g, Functor h, f :<: g) => f :<: (h :+: g) where
  inj = R . inj
```

Using this type class, smart constructors can be defined.

```
inject :: (g :<: f) => g (Fix f) -> Fix f
inject = In . inj
val :: (ValF2 :<: f) => Int -> Fix f
```

```

val x = inject (ValF2 x)
mul :: (MulF2 :-: f) ⇒ Fix f → Fix f → Fix f
mul x y = inject (MulF2 x y)

```

Expressions can now be built using the constructors, such as `val 12 `mul` val 34`.

A modular algebra can now be defined that provides an interpretation of this datatype.

```

class Functor f ⇒ EvalAlg f where
  evalAlg :: f Int → Int
instance (EvalAlg f, EvalAlg g) ⇒ EvalAlg (f :+: g) where
  evalAlg (L x) = evalAlg x
  evalAlg (R y) = evalAlg y
instance EvalAlg MulF2 where
  evalAlg (MulF2 x y) = x * y
instance EvalAlg ValF2 where
  evalAlg (ValF2 x) = x
cata :: Functor f ⇒ (f a → a) → Fix f → a
cata alg (In x) = alg (fmap (cata alg) x)
eval :: EvalAlg f ⇒ Fix f → Int
eval = cata evalAlg

```

One benefit to this approach is that an interpretation is only needed for expressions that only use `MulF` and `ValF`. If a new constructor such as `SubF` was added to the language and it would never be given to this fold, then it would not require an instance. This helps to solve the expression problem.

2.5 Dependently Typed Programming

Although Haskell does not officially support dependently typed programming, there are techniques available that together can be used to replicate some of the experience.

2.5.1 DataKinds Language Extension

Through the use of the `DataKinds` language extension [27], all data types can be promoted to also be kinds and their constructors to be type constructors. When constructors are promoted to type constructors, they are prefixed with a `'`. This allows for more interesting and restrictive types.

Consider the example of a vector that also maintains its length. Peano numbers can be used to keep track of the length, which prevents a negative length for a vector. This is where numbers are defined as zero or a number `n` incremented by 1.

```

data Nat = Zero
         | Succ Nat

```

A vector type can now be defined that makes use of the promoted `Nat` kind.

```

data Vec :: Type → Nat → Type where
  Nil    :: Vec a 'Zero
  Cons :: a → Vec a n → Vec a ('Succ n)

```

The use of `DataKinds` can enforce stronger types. For example a function can now require that a specific length of vector is given as an argument. With standard lists, this would not be possible, which could result in run-time errors when the incorrect length is used. For example, getting the head of a list. Getting the head of an empty list an error will be thrown. For a vector, a `safeHead` function can be defined that will not type check if the vector is empty.

```

safeHead :: Vec a ('Succ n) → a
safeHead (Cons x _) = x

```

2.5.2 Singletons

DataKinds are useful for adding extra information back into the types, but how can information be recovered from the types? For example, could a function that gets the length of a vector be defined?

```
vecLength :: Vec a n → Nat
```

This is enabled through the use of singletons [6]. A singleton in Haskell is a type that has just one inhabitant. That is that there is only one possible value for each type. They are written in such a way that pattern matching reveals the type parameter. For example, the corresponding singleton instance for `Nat` is `SNat`. The structure for `SNat` closely flows that of `Nat`.

```
data SNat (n :: Nat) where
  SZero :: SNat 'Zero
  SSucc :: SNat n → SNat ('Succ n)
```

A function that fetches the length of a vector can now definable.

```
vecLength2 :: Vec a n → SNat n
vecLength2 Nil          = SZero
vecLength2 (Cons x xs) = SSucc (vecLength2 xs)
```

2.5.3 Type Families

Now consider the possible scenario of appending two vectors together. How would the type signature look? This leads to the problem where two type-level `Nats` need to be added together. This is where Type Families [18] become useful, they allow for the definition of functions on types. Consider the example of appending two vectors together, this would require type-level arithmetic — adding the lengths together.

```
vecAppend :: Vec a n → Vec a m → Vec a (n :+ m)
```

This requires a `:+` type family that can add two `Nats` together.

```
type family (a :: Nat) :+ (b :: Nat) where
  a :+'Zero    = a
  a :+'Succ b = 'Succ (a :+ b)
```

2.5.4 Summary

Together these features allow for dependently typed programming constructs in Haskell:

- DataKinds allow for values to be promoted to types
- Singletons allow types to be demoted to values
- Type Families can be used to define functions that manipulate types.

2.6 Heterogeneous Lists

```
data HList (xs :: [Type]) where
  HCons :: x → HList xs → HList (x' : xs)
  HNil  :: HList '[]
```

Add an example of simple `HList`

Add an example to get the length of a type list. Length needed later

2.7 Existential Types

Existential Types

2.8 Phantom Type Parameters

Phantom type parameters

2.9 Monadic Resource Theories

Monadic resource theories — Symmetric Monoidal pre-orders.

Chapter 3

The Language

3.1 Language Requirements

For the design of the language to be considered a success, several criteria need to be met:

- **Describe any dataflow** — The combinators in the language should be able to describe any dataflow that the user needs.
- **Quickness to Learn** — The language should be quick and easy to learn. If it takes users too long to learn, they may never bother, meaning that the language will never be used.
- **Easy to write** — A user should be able to write programs easily. They should not have to spend time creating additional boilerplate code, where it is not necessary.
- **Easy to understand** — Any program written with this language should be easy to understand, so that users can review existing code and know what it will do.
- **Type-safe** — A feature missing in many dataflow tools is the lack of type checking. This causes problems later on in the development process with more debugging and testing needed. The language should be type-safe to avoid any run-time errors occurring where types do not match.

3.2 Tasks

Dataflow programming focuses on the transforming inputs to outputs. To be able to transform inputs this language will make use of tasks. They are responsible for reading from an input data source, completing some operation on the input, then finally writing to an output data sink. Tasks could take many different forms, for example they could be:

- A pure function — a function with type $a \rightarrow b$
- An external operation — interacting with some external system. For example, calling a terminal command.

A task could have a single input or multiple inputs, however, for now just a single input task will be considered. Multi-input tasks are explained further in Sub-Section [3.4.3](#)

3.2.1 Data Stores

Data stores are used to pass values between different tasks, this ensures that the input and output of tasks are closely controlled. They are used to represent different ways of storing values: one example could be a point to a CSV file. A data store can be defined as a type class, with two methods `fetch` and `save`:

Motivate further why a `DataStore` needs to exist — prevents the user from reading from a source incorrectly.

```

class DataStore f a where
  fetch :: f a → IO a
  save  :: f a → a → IO (f a)

```

A `DataStore` has two type parameters: where `f` is the type of `DataStore` being used and `a` is the type of the value stored inside it. The aptly named methods describe their intended function: `fetch` will fetch a value from a `DataStore`, and `save` will save a value. The `fetch` method takes a `DataStore` as input and will return the value stores inside. However, the `save` method may not be as self explanatory, since it has an extra `f a` argument. This argument can be thought of as a pointer to a `DataStore`: it contains the information needed to save. For example, in the case of a file store it could be the file name.

By implementing this as a type class, there can be many different implementations of a `DataStore`. The library comes with several pre-defined `DataStores`, such as a `VariableStore`. This can be thought of as an in memory storage between tasks.

```

data VariableStore a = Var a | Empty
instance DataStore VariableStore a where
  fetch (Var x) = return x
  save Empty x = return (Var x)

```

The `VariableStore` is the most basic example of a `DataStore`, a more complex example is a `FileStore`, which represents a pointer to a file:

```

newtype FileStore a = FileStore String
instance DataStore FileStore String where
  fetch (FileStore fname) = readFile fname
  save (FileStore fname) x = writeFile fname x >> return (FileStore fname)
instance DataStore FileStore [String] where
  fetch (FileStore fname) = readFile fname >> return · lines
  save (FileStore fname) x = writeFile fname (unlines x) >> return (FileStore fname)

```

The `FileStore` is only defined to store two different types: `String` and `[String]`. If a user attempts to store anything other than these two types then a compiler error will be thrown, for example:

```

ghci> save (FileStore "test.txt") (123 :: Int)
> No instance for (DataStore FileStore Int) arising from a use of ‘save’

```

Although a small set of `DataStores` are included in the library, the user is also able to add new instances of the type class with their own `DataStores`. Some example expansions, could be supporting writing to a database table, or a Hadoop file system.

3.2.2 Task Constructor

The type of a task details the inputs and outputs. A task is created via a constructor that takes two arguments: the function it represents and somewhere to store the output. This constructor makes use of GADTs syntax [17] so that constraints can be placed on the types used. It enforces that a `DataStore` must exist for the input and output types. This allows the task to make use of the `fetch` and `save` functions.

```

data Task (f :: Type → Type) (a :: Type) (g :: Type → Type) (b :: Type) where
  Task :: (DataStore f a, DataStore g b) ⇒ (f a → g b → IO (g b)) → g b → Task f a g b

```

When a `Task` is executed the stored function is executed, with the input being passed in as the first argument and the output “pointer” as the second argument. This returns an output `DataStore` that can be passed on to another `Task`.

3.3 First Attempt: Chains

In a dataflow programming, one of the key aspects is the definition of dependencies between tasks in the flow. One possible approach to encoding this concept in the language is to make use of sequences of tasks — also referred to as chains. These chains compose tasks, based on their dependencies. A chain can be modelled with an abstract datatype:

```

data Chain (f :: Type → Type) (a :: Type) (g :: Type → Type) (b :: Type) where
  Chain :: Task f a g b → Chain f a g b
  Then :: Chain f a g b → Chain g b h c → Chain f a h c

```

free moniod? maybe yes it looks like what eddie did...

The Chain constructor allows for a Task to be wrapped in the Chain data type. This allows for them to be easily composed with other Tasks. Without this there would need to be an “empty” element, however, this could lead to the construction of a chain with no tasks. The Then constructor is used to combine multiple chains to form a sequence of sequential tasks. To make this easier to use an operator >>> is defined that represents Then:

```

(>>>) :: Chain f a g b → Chain g b h c → Chain f a h c
(>>>) = Then

```

This can now be used to join sequences of tasks together, for example:

```

task1 :: Task VariableStore Int      VariableStore String
task2 :: Task VariableStore String  FileStore      [String]
task3 :: Task FileStore      [String] VariableStore Int
sequence :: Chain VariableStore Int VariableStore Int
sequence = Chain task1 >>> Chain task2 >>> Chain task3

```

sequence will perform the three tasks in order, starting with task1 and finishing with task3.

3.3.1 Trees as Chains

Now that tasks can be performed in sequence, the next logical step will be to introduce the concept of branching out. This results in a task's output being given to multiple tasks, rather than just 1.

To do this a new abstract datatype is required. This will be used to form a list of Chains, conventionally the [] type would be used, however this is not possible as each chain will have a different type. This means that existential types will need to be used.

```

data Pipe where
  Pipe :: ∀ f a g b. (DataStore f a, DataStore g b) ⇒ Chain f a g b → Pipe
  And :: Pipe → Pipe → Pipe

```

cite where
this comes
from?

The And constructor can be used to combine multiple chains together. Figure 3.1 shows the previous sequence, with a new task4 which also uses the input from task2.

However, there is a problem with this approach: to be able to form a network similar to that shown in Figure 3.1, the language will need to know where to join two Chains together. However, with the current definition of a Task, it is not possible to easily check the equivalence of two functions. Being able to check for equivalence will be key to defining a method to merge these chains together: it will need to match task2 in the first chain with the task2 in the second chain. Similarly, if a user wanted to use the same task multiple times, it would not be possible to differentiate between them. One approach to this would be to have unique identifiers for each task, such as PIDs.

Process Identifiers (PIDs) A Chain can be modified so that instead of storing a Task it instead stores a PID — a unique identifier for a task. However to do this a new PID data type is needed:

```

data PID (f :: Type → Type) (a :: Type) (g :: Type → Type) (b :: Type) where
  PID :: Int → PID f a g b

```

The PID data type has the same kind as a Task and makes use of phantom type parameters, to retain the same type information as a Task, whilst storing just an Int that can be used to identify it.

```

data Chain' (f :: Type → Type) (a :: Type) (g :: Type → Type) (b :: Type) where
  Chain' :: PID f a g b → Chain' f a g b
  Then' :: Chain' f a g b → Chain' g b h c → Chain' f a h c

```


This new version of Chain named Chain', is almost identical in the way it behaves, however instead of storing a Task, it stores a PID.

```

task4 :: Task FileStore [String] VariableStore String
branchExample :: Pipe
branchExample = Pipe (Chain task1 >>> Chain task2 >>> Chain task3)
                `And`
                Pipe (Chain task2 >>> Chain task4)

```

(a)



(b)

Figure 3.1: A Pipe (a) and its corresponding dataflow diagram (b).

This, however, leaves a key question, how do **Tasks** get mapped to **PIDs**. This can be done by employing the **State** monad. The state stores a map from **PID** to task and a counter for **PIDs**. As **Tasks** each have a different type a new datatype is required to use existential types, to close over the types, and produce values of the same type. This is because a **Map** can only store one type. The **Workflow** monad is a type alias for the **State** monad, which stores the **WorkflowState**.

```

data TaskWrap = ∀ f a g b. TaskWrap (Task f a g b)
data WorkflowState = WorkflowState {
  pidCounter :: Int,
  tasks :: M.Map Int TaskWrap
}
type Workflow = State WorkflowState

```

An operation that is defined in the monad is **registerTask**. This takes a **Task** and returns a **Chain** that stores a **PID** inside it. Whenever a user would like to add a new task to the workflow, they register it. They can then use this returned value to construct multiple chains, which can now be joined easily by comparing the stored **PID**.

```
registerTask :: Task f a g b → Workflow (Chain' f a g b)
```

One benefit to this approach is that if the user would like to use a task again in a different place, they can simply register it again and use the new **PID** value.

3.3.2 Evaluation

☑ **Describe any Dataflow** This method is capable of describing any dataflow, although in its current state, it can only support trees. The algorithm that would be used to join different chains together could be developed to allow them to rejoin onto another chain. This will allow for any Directed Acyclic Graph (**DAG**) to be defined.

☑ **Quickness to Learn** With only 4 different combinators and the **Task** constructor, this language is very simple to learn. There are only two main concepts the user needs to understand: how to join tasks into a chain and how to join chains together.

☑ **Easy to Write** Building chains is a very simple process, and allows the user to focus on the dependencies of one task. They do not need to be aware of the bigger picture. For example, adding a new **task5** that depends on **task2** in Figure 3.1. The user will just need to add one extra chain **task5 >>> task2**, which will have no effect on the existing chains.

☒ **Easy to Read** Although it is easy to write chains, this could lead to messy definitions with no structure. This will make it harder for a user to interpret an already defined collection of chains. Its possible that they will need to draw out some of the dependencies to further understand what is already defined.

☒ **Type-safe** Although a `Chain` can be well typed, the use of existential types to join chains together pose a problem. This causes the types to be ‘hidden’, this means that when executing these tasks, the types need to be recovered. This is possible through the use of `gcast` from the `Data.Typeable` library. However, this has to perform a reflection at run-time to compare the types. There is the possibility that the types could not matching and this would only be discovered at run-time. There is a mechanism to handle the failed match case, however, this does not fulfil the criteria of being fully type-safe.

Chains are not good enough Chains do not satisfy some of the key requirements that this library set out to solve, such as type safety. This could lead to crashes at run-time as the compiler is not able to fully verify that the system type-checks. Another approach could be to look to category theory, for ways to compose functions together in a type-safe way.

3.4 Final Attempt: Circuit

This approach was inspired by the parallel prefix circuits as described by Hinze [9]. It uses constructors similar to those used by Hinze to create a circuit that represents the **DAG**, used in the dataflow. The constructors seen in Figure 3.2 represent the behaviour of edges in a graph.

Opinions on way to draw `DropL/R??` or do i change it to `Drop`, which aligns with MRT.

maybe this is something to do with MRT??




Figure 3.2: The constructors in the `Circuit` library alongside their graphical representation.

3.4.1 Constructors

Each of these constructors use strong types to ensure that they are combined correctly. A `Circuit` has 7 different type parameters:

```
Circuit (inputsStorageTypes :: [Type → Type]) (inputsTypes :: [Type]) (inputsApplied :: [Type])
        (outputsStorageTypes :: [Type → Type]) (outputsTypes :: [Type]) (outputsApplied :: [Type])
        (nInputs :: Nat)
```

A `Circuit` can be thought of as a list of inputs, which are processed and a resulting list of outputs are produced. To represent this it makes use of the `DataKinds` language extension [27], to use type-level lists and natural numbers. Each parameter represents a certain piece of information needed to construct a circuit:

- `inputsStorageTypes` is a type-list of storage types, for example `'[VariableStore, CSVStore]` — these all have kind $* \rightarrow *$.
- `inputsTypes` is a type-list of the types stored in the storage, for example `'[Int, [(String, Float)]]`.
- `inputsApplied` is a type-list of the storage types applied to the types stored, for example `'[VariableStore Int, CSVStore [(String, Float)]]`.
- `outputsStorageTypes`, `outputsTypes` and `outputsApplied` mirror the examples above, but for the outputs instead.
- `nInputs` is a type-level `Nat` that is the length of the input lists.

In the language there are two different types of constructor, those that recurse and those that can be considered leaf nodes. The behaviour of both types of constructor is recorded within the types. For example, the `id` constructor has the type:

```
id :: DataStore' '[f] '[a] ⇒ Circuit '[f] '[a] '[f a] '[f] '[a] '[f a] N1
```

It can be seen how the type information for this constructor states that it has 1 input value of type `f a` and it returns that same value. Each type parameter in `id` is a phantom type [3], since there are no values stored in the data type that use the type parameters. The rest of the constructors that are leaf nodes are: `replicate`, `swap`, `dropL`, and `dropR`:

```
replicate :: DataStore' '[f] '[a] ⇒ Circuit '[f] '[a] '[f a] '[f, f] '[a, a] '[f a, f a] N1
swap      :: DataStore' '[f, g] '[a, b] ⇒ Circuit '[f, g] '[a, b] '[f a, g b] '[g, f] '[b, a] '[g b, f a] N2
dropL     :: DataStore' '[f, g] '[a, b] ⇒ Circuit '[f, g] '[a, b] '[f a, g b] '[g] '[b] '[g b] N2
dropR     :: DataStore' '[f, g] '[a, b] ⇒ Circuit '[f, g] '[a, b] '[f a, g b] '[f] '[a] '[f a] N2
```

The `replicate` constructor states that a single input value of type `f a` should be input, and that value should then be duplicated and output. The `swap` constructor takes two values as input: `f a` and `g b`. It will then swap these values over, such that the output will now be: `g b` and `f a`. `dropL` will take two inputs: `f a` and `g b`. It will then drop the left argument and return just a `g b`. The `dropR` has the same behaviour as `dropL`, it just drops the right argument instead.

To be able to make use of the leaf nodes, they need to be combined in some way. To do this two new constructors named ‘`beside`’ and ‘`then`’ will be used. However, before defining these constructors there are some tools that are required. This is due to the types no longer being concrete. For example, the input type list is no longer known: it can only be referred to as `fs` and `as`. This means it is much harder to specify the new type of the `Circuit`.

Apply Type Family It would not be possible to use a new type variable `xs` for the `inputsApplied` parameter. This is because it needs to be constrained so that it is equivalent to `fs` applied to `as`. To solve this a new closed type family [5] is created that is able to apply the two type lists together. This type family pairwise applies a list of types with kind $* \rightarrow *$ to a list of types with kind $*$ to form a new list containing types of kind $*$. For example, `Apply '[f, g, h] '[a, b, c] ~ '[f a, g b, h c]`.

': looks awful

```
type family Apply (fs :: [Type → Type]) (as :: [Type]) where
  Apply '[] '[] = '[]
  Apply (f' : fs) (a' : as) = f a' : Apply fs as
```

Append Type Family There will also be the need to append two type level lists together. For example, `'[a, b, c] :++' [d, e, f] ~ '[a, b, c, d, e, f]`. To do this an append type family [12] can be used:

```
type family (:++) (l1 :: [k]) (l2 :: [k]) :: [k] where
  (:++) '[] l = l
  (:++) (e' : l) l' = e' : (l :++ l')
```

The `:++` type family is defined in the same way as the standard `++` function on value lists, however, it appends type lists together instead. This type family makes use of the language extension `PolyKinds` [27] to allow for the append to be polymorphic on the kind stored in the type list. This will avoid defining multiple versions to append `fs` with `gs`, and `as` with `bs`.

The ‘Then’ Constructor This constructor — denoted by $\langle\Rightarrow\rangle$ — is used to stack two circuits on top of each other. It is used to encapsulate the idea of dependencies, between different circuits. Through types it enforces that the output of the top circuit is the same as the input to the bottom circuit.

```
(⟨⇒⟩) :: (DataStore' fs as, DataStore' gs bs, DataStore' hs cs)
      ⇒ Circuit fs as (Apply fs as) gs bs (Apply gs bs) nfs
      → Circuit gs bs (Apply gs bs) hs cs (Apply hs cs) ngs
      → Circuit fs as (Apply fs as) hs cs (Apply hs cs) nfs
```

It employs a similar logic to function composition $(\cdot) :: (a \rightarrow b) \rightarrow (b \rightarrow c) \rightarrow (a \rightarrow c)$. The resulting type from this constructor uses the input types from the first argument $fs\ as\ (Apply\ fs\ as)$, and the output types from the second argument $hs\ cs\ (Apply\ hs\ cs)$. It then forces the constrain that the output type of the first argument and the input type of the second are the same — $gs\ bs\ (Apply\ gs\ bs)$.

The ‘Beside’ Constructor Denoted by $\langle\>$, the beside constructor is used to place two circuits side-by-side. The resulting `Circuit` has the types of left and right circuits appended together.

```
(⟨⟩) :: (DataStore' fs as, DataStore' gs bs, DataStore' hs cs, DataStore' is ds)
      ⇒ Circuit fs as (Apply fs as) gs bs (Apply gs bs) nfs
      → Circuit hs cs (Apply hs cs) is ds (Apply is ds) nhs
      → Circuit (fs :++ hs) (as :++ cs) (Apply fs as :++ Apply hs cs)
                (gs :++ is) (bs :++ ds) (Apply gs bs :++ Apply is ds)
                (nfs :++ nhs)
```

This constructor works by making use of the $:++$ type family to append the input and output type list of the left constructor to those of the right constructor. It also makes use of the $:+$ type family — defined in Section ?? — to add the number of inputs from the left and right together.

The $\langle\>$ constructor combined with the `id` constructor can be considered a Monoid at the behaviour level. However, due to the types on the constructors it is not possible to define a Monoid instance for $\langle\>$ in Haskell.

Give example on the types as to why it isnt.

I think its a monoid in cat theory, but it isn't one in haskell. empty exists with id, but this would not conform to the types.

3.4.2 Combined Data Stores

`DataStore'` does not need `save'` as it is never used in a task. Can only have 1 output.

A keen eyed reader may notice that all of these constructors have not been using the original `DataStore` type class. Instead they have all used the `DataStore'` type class. This is a special case of a `DataStore`, allowing for constructors to also be defined over type lists, not just a single type. Combined data stores make it easier for tasks to fetch from multiple inputs. Users will just have to call a single `fetch'` function, rather than multiple. However, since tasks can only have one output, there is no need for a `save'` function, that would be able to save to multiple data stores.

To be able to define `DataStore'`, heterogeneous lists [12] are needed — specifically two different forms. `HList'` stores values of type `f a` and is parameterised by two type lists `fs` and `as`. `IOList` stores items of type `IO a` and is parameterised by a type list `as`. Using an `IOList` makes it easier to define a function that produces a list of IO computations. Their definitions are:

```
data HList' (fs :: [Type → Type]) (as :: [Type]) where
  HCons' :: f a → HList' fs as → HList' (f' : fs) (a' : as)
  HNil'  :: HList' '[] '[]

data IOList (xs :: [Type]) where
  IOCons :: IO x → IOList xs → IOList (x' : xs)
  IONil  :: IOList '[]
```

Now that there is a mechanism to represent a list of different types, it is possible to define `DataStore'`:

```
class DataStore' (fs :: [Type → Type]) (as :: [Type]) where
  fetch' :: HList' fs as → IOList as
```

To save the user of the cumbersome task of having to define an instance of `DataStore'` for every possible combination of data stores, the instance is derived from the previous `DataStore` type class. This means that a user does not need to create any instances of `DataStore'`. They can instead focus on each single case, with the knowledge that they will automatically be able to combine them with other data stores.

```
instance {-# OVERLAPPING #-} (DataStore f a) => DataStore' '[f]' '[a] where
    fetch' (HCons' x HNil') = IOCons (fetch x) IONil
instance (DataStore f a, DataStore' fs as) => DataStore' (f' : fs) (a' : as) where
    fetch' (HCons' x xs) = IOCons (fetch x) (fetch' xs)
```

One of these instances makes use of the `{-# OVERLAPPING #-}` pragma. In most cases the base case instance that would be defined is `DataStore' '[]' '[]`. However, it does not make sense to have an empty data store. Therefore, the base case is selected to be a list with one element `DataStore' '[f]' '[a]`. This leads to a problem: GHC is unable to decide which instance to use. It could use either the `DataStore' '[f]' '[a]` or the `DataStore' (f' : []) (a' : [])` instance. The overlapping pragma tells GHC, that if it encounters this scenario, it should choose the one with the pragma.

3.4.3 Multi-Input Tasks

With a `Circuit`, it is possible to represent a **DAG**. This means that a node in the graph can now have multiple dependencies, as seen in Figure 3.3.




Figure 3.3: A graphical representation of a task with multiple dependencies

To support this, a modification is made to the `task` constructor: rather than have an input value type of `f a`, it can now have an input value type of `HList' fs as`. The function executed in the task can now use `fetch'` to fetch all inputs with one function call.

```
task :: (DataStore' fs as, DataStore g b)
    => (HList' fs as -> g b -> IO (g b))
    -> g b
    -> Circuit fs as (Apply fs as) '[g]' '[b]' '[g b]' (Length fs)
```

Now that the length of the inputs is unknown, in order to specify the `nInputs` type parameter, the `Length` type family defined in Section 2.6 must be used. This will return a type-level `Nat`, which is the length of the input array `fs`.

Smart Constructors There could be many times that the flexibility provided by defining your own tasks from scratch could cause a large amount of boiler plate code. For example, there may be times that a user already has pre-defined function and would like to convert it to a task. Therefore there are also two smart constructors that they are able to use:

```
multiInputTask :: (DataStore' fs as, DataStore g b)
    => (HList as -> b)
    -> g b
    -> Circuit fs as (Apply fs as) '[g]' '[b]' '[g b]' (Length fs)
functionTask :: (DataStore f a, DataStore g b)
    => (a -> b)
    -> g b
    -> Circuit '[f]' '[a]' '[f a]' '[g]' '[b]' '[g b]' N1
```

The first allows for a simple function with multiple inputs to be defined. With the fetching and saving handled by the smart constructor. The second allows for a simple `a -> b` function to be turned into a `Task`.

3.4.4 mapC operator

Currently a circuit has a static design — once it has been created it cannot change. There are times when this could be a flaw in the language. For example, when there is a dynamic number of inputs. This could be combated with more smart constructors to generate more complex circuits, with the pre-existing constructors. Another approach would be to add new constructors that allow for more dynamic circuits, such as `mapC`. This new constructor is used to map a circuit on a single input containing a list of items. The input is fed into the inner circuit, accumulated back into a list, and then output.

```
mapC :: (DataStore' '[f] '[[a]], DataStore g [b])
  => Circuit '[VariableStore] '[a] '[VariableStore a] '[VariableStore] '[b] '[VariableStore b] N1
  -> g [b]
  -> Circuit '[f] '[[a]] '[f [a]] '[g] '[[b]] '[g [b]] N1
```

Graphical representation of this?

3.4.5 Completeness

something about the stuff Alex said

monadic resource theories.

3.4.6 Evaluation

☑ **Describe any Dataflow** Due to the use of constructors which align with a symmetric monoidal pre-order, it can be concluded that it is possible to define any **DAG** needed for a dataflow.

I can't find anything to cite that discusses this though :(, but afaik i cannot find a scenario that cannot be made...

☑ **Quickness to Learn** The language uses, combinators that also have a visual representation, this makes it easy to quickly understand how they all work. A user can then also benefit from the familiarity of using the host language Haskell.

☐ **Easy to Write** Building a circuit is a more difficult task: to be able to define it the user needs to have an understanding of the shape of the dataflow diagram. One the user has a sketch for the dataflow they would like to create, translating to a circuit is a much simpler job. This creates more upfront work for the user, however, it is offset by the additional benefits that a circuit brings.

☑ **Easy to Read** Due to the graphical nature of the constructors, it is relatively simple to build up a picture of how an existing circuit works. The user is able to infer the shape of the dataflow diagram easily by working their way down a circuit from top to bottom, and visualising how different tasks are connected. It could also be possible to pretty print a circuit to recreate the dataflow diagram — although this has not been implemented.

☑ **Type-safe** One key benefit that a `Circuit` brings is that constructing them uses strong types. Each constructor encodes its behaviour within the types. This allows the GHC type checker to validate a `Circuit` at compile-time, to ensure that each task is receiving the correct values. This avoids the possibility of crashes at run-time, where types do not match correctly. There is, however, a consequence of this type-safety: the user now needs to add some explicit types on a `Circuit` to help the type checker.

Easy to Build A `Circuit` focuses on the transformations that are made to edges on a graph. This can be beneficial to the user as it is the edges in a dataflow diagram that encode dependencies between tasks. Although circuits may initially appear complex, there is a relatively simple process that can be used to construct them. By hand-drawing a dataflow diagram, a circuit can always be constructed that closely mirrors this diagram. This means that the user can easily visualise what is happening inside a circuit.

Chapter 4

Implementation

4.1 Requirements

The implementation of the network itself has several requirements that are separate from the API that the user will use:

- **Type-safe** — This is a continuation of the previous requirement for the language. It is also important that once the user has built a well-typed `Circuit`, that the code also continues to be executed in a well-typed environment, to ensure that all inputs and outputs are correctly typed.
- **Parallel** — One of the key benefits that comes from dataflow programming is implicit parallelisation. With this **DSL** being tailored towards data pipelines, which could be computationally expensive, it should be able to benefit from parallel execution.
- **Competitive Speed** — This library should be able to execute dataflows in a competitive time, with other libraries that already exist.
- **Failure Tolerance** — It is important that if one invocation of a task crashes, it does not crash the whole program. This implementation should be able to gracefully handle errors and propagate them through the circuit.

4.2 Circuit AST

The constructors for the language are actually *smart constructors*. They provide a more elegant way to build an **AST**, which represents the circuit.

flesh this out a bit more, doesnt really fit in the current context

??? Is there anything else that maybe I haven't achieved? memorisation?

4.2.1 IFunctor

The **AST** that represents a `Circuit` is built using indexed functors, also known as `IFunctor`. `IFunctors` are, however, only defined to take a single type index — a `Circuit` needs 7. To do this `IFunctor7` can be defined:

```
class IFunctor7 iF where
  imap7 :: (∀ a b c d e f g. f' a b c d e f g → g' a b c d e f g) → iF f' a b c d e f g → iF g' a b c d e f g
```

`IFunctor7` follows a similar structure to a standard `IFunctor`, it just has 7 type indices instead. This indexed functor can be used to mark the recursive points of the data types used to construct the **Circuit AST**.

However, any data type that is converted to use an `IFunctor7`, will need a way to tie the recursive knot. A new type called `IFix7` can be used, it will follow a similar pattern to `IFix`, but with 7 type indices.

```
newtype IFix7 iF a b c d e f g = lln7 (iF (IFix7 iF) a b c d e f g)
```

4.2.2 Type-Indexed Data Types à la Carte

To build the **AST**, the data types à la carte approach is taken. This allows for a modular approach, making the library more extendable later on. To be able to use this approach, it needs to be modified to support indexed functors. This modification, changes them to take an indexed functor with 7 type indices — a through to g:

cite

```

data (iF :: iG) (f' :: i → j → k → l → m → n → o → Type)
  (a :: i) (b :: j) (c :: k) (d :: l) (e :: m) (f :: n) (g :: o) where
  L :: iF f' a b c d e f g → (iF :: iG) f' a b c d e f g
  R :: iG f' a b c d e f g → (iF :: iG) f' a b c d e f g
infixr :+:

```

As mentioned in Section 2.4, using the :+: operator comes with problem of many L's and R's, when trying to create the **AST** structure. To avoid this, the :-< operator can also be extended to work with IFunctor₇. The definition follows closely to the original definition, with some minor modifications to the types and constraints.

```

class (IFunctor7 iF, IFunctor7 iG) ⇒ iF :-< iG where
  inj :: iF f' a b c d e f g → iG f' a b c d e f g
instance IFunctor7 iF ⇒ iF :-< iF where
  inj = id
instance (IFunctor7 iF, IFunctor7 iG) ⇒ iF :-< (iF :+: iG) where
  inj = L
instance (IFunctor7 iF, IFunctor7 iG, IFunctor7 iH, iF :-< iG) ⇒ iF :-< (iH :+: iG) where
  inj = R · inj

```

Defining a constructor Data types for each constructor can be defined individually. The Then constructor is used as an example, however, the process can be applied to all constructors in the language.

```

data Then (iF :: [Type → Type] → [Type] → [Type]
  → [Type → Type] → [Type] → [Type] → Nat → Type)
  (inputsS :: [Type → Type]) (inputsT :: [Type]) (inputsA :: [Type])
  (outputsS :: [Type → Type]) (outputsT :: [Type]) (outputsA :: [Type])
  (ninputs :: Nat) where
  Then :: (DataStore' fs as, DataStore' gs bs, DataStore' hs cs)
    ⇒ iF fs as (Apply fs as) gs bs (Apply gs bs) nfs
    → iF gs bs (Apply gs bs) hs cs (Apply hs cs) ngs
    → Then iF fs as (Apply fs as) hs cs (Apply hs cs) nfs

```

Each iF denotes the recursive points in the data type, with the subsequent type arguments mirroring those seen in Section 3.4.1. A corresponding IFunctor₇ instance formalises the points of recursion, by showing how to transform the structure inside it.

```

instance IFunctor7 Then where
  imap7 f (Then x y) = Then (f x) (f y)

```

The smart constructor, that injects the L's and R's automatically can be defined for Then as:

```

(<=>) :: (Then :-< iF, DataStore' fs as, DataStore' gs bs, DataStore' hs cs)
  ⇒ IFix7 iF fs as (Apply fs as) gs bs (Apply gs bs) nfs
  → IFix7 iF gs bs (Apply gs bs) hs cs (Apply hs cs) nhs
  → IFix7 iF fs as (Apply fs as) hs cs (Apply hs cs) nfs
(<=>) l r = lln7 (inj (Then l r))
infixr 4 <=>

```

The constructor adds one extra constraint, to the constructor defined in Section 3.4.1 — Then :-< iF. This allows the smart constructor to produce an node in the **AST** for any co-product of data types, that includes the Then data type.

Representing a Circuit Once each constructor has been defined then they can be combined together to form the `CircuitF` type, which can be used to represent a circuit.

```
type CircuitF = Id :+: Replicate :+: Then :+: Beside :+: Swap :+: DropL :+: DropR :+: Task :+: Map
```

combined
= co-
product

The fixed-point of the `CircuitF` datatype can be defined with `IFix7`:

```
type Circuit = IFix7 CircuitF
```

4.3 Process Network

Now that it is possible to build a `Circuit`, which can be considered a specification for how to execute a set of tasks, there needs to be a mechanism in place to execute the specification. The standard implementation of a process network will use a Kahn Process Network (**KPN**). This means that each task in a circuit will run on its own separate thread, with inputs being passed between them on unbounded channels.

4.3.1 Network Typeclass

To allow for different process networks, a typeclass will be used to specify all the functions that every network should have. The `Network` typeclass is defined as:

```
class Network n where
  startNetwork :: Circuit inputsS inputsT inputsA outputsS outputsT outputsA nInputs
               → IO (n inputsS inputsT inputsA outputsS outputsT outputsA)
  stopNetwork  :: n inputsS inputsT inputsA outputsS outputsT outputsA
               → IO ()
  write :: HList' inputsS inputsT
        → n inputsS inputsT inputsA outputsS outputsT outputsA
        → IO ()
  read  :: n inputsS inputsT inputsA outputsS outputsT outputsA
        → IO (HList' outputsS outputsT)
```

This type class requires that a network has 4 different functions:

- `startNetwork` is responsible for converting the circuit into the underlying representation for a process network: it will be discussed in more detail in Section 4.4.
- `stopNetwork` is for cleaning up the network after it is no longer needed. For example, this could be stopping the threads running. This could be particularly important if embedding a circuit into a larger program, where unused threads could be left hanging.
- `write` should take some input values and add them into the network, so that they can be processed.
- `read` should retrieve some output values from the network.

Interaction with Network

Should I add an example of how to use it?

4.3.2 The Basic Network Representation

An implementation of the `Network` typeclass is a `BasicNetwork`. This implementation makes use of a special case of heterogeneous list. A `PipeList` is used to represent a heterogeneous list of channels. This allows the `BasicNetwork` to store multiple channels in the same list without the need for existential types — one of the problems previously encountered with chains in Section 3.3.

```
data PipeList (fs :: [Type → Type]) (as :: [Type]) (xs :: [Type]) where
  PipeCons :: Chan (fa) → PipeList fs as xs → PipeList (f' : fs) (a' : as) (fa' : xs)
  PipeNil  :: PipeList '[] '[] '[]
```

Does basic make it sound simple? should i got with standard or something like that??

Making use of `PipeLists`, the `BasicNetwork` data type can be defined. This definition makes use of record syntax, this allows for named fields, with accessors automatically generated.

```
data BasicNetwork (inputsS :: [Type → Type]) (inputsT :: [Type]) (inputsA :: [Type])
  (outputsS :: [Type → Type]) (outputsT :: [Type]) (outputsA :: [Type]) where
  BasicNetwork :: {
    threads :: [ThreadId],
    inputs  :: PipeList inputsS inputsT inputsA,
    outputs :: PipeList outputsS outputsT outputsA }
    → BasicNetwork inputsS inputsT inputsA outputsS outputsT outputsA
```

The `BasicNetwork` has three fields:

- `threads` is a list of `ThreadId`s, this allows for the threads to be managed after their creation.
- `inputs` is a `PipeList` containing the channels that the initial input into the network.
- `outputs` is a `PipeList` that stores channels, which the output of the network can be read from.

The `Network` type instance for a `BasicNetwork` is relatively trivial to implement: if given a function to transform a `Circuit` to it.

```
instance Network BasicNetwork where
  startNetwork = buildBasicNetwork -- Definition to come...
  stopNetwork n = forM_ (threads n) killThread
  write uuid xs n = writePipes xs (inputs n)
  read n         = readPipes (outputs n)
```

The `writePipes` function will input a list of values into each of the respective pipes. The `readPipes` function will make a blocking call to each channel to read an output from it. This function will block till an output is read from every output channel.

Should I add the definition of these?

4.4 Translation to a Network

There is now a representation for a `Circuit` that the user will build, and a representation used to execute the `Circuit`. However, there is no mechanism to convert between them. This can be achieved by folding the circuit data type into a network. This fold, however, will need to create threads and channels, both of which IO actions. The current definition for the fold `icata7` is not able perform monadic computation inside the algebra. To solve this `unsafePerformIO` could be used, however, for this to be safe the IO computation needs to have no side-effects. This fold will violate this rule, therefore, the only other way to support this is to modify the catamorphism to support monadic computation.

4.4.1 Indexed Monadic Catamorphism — `icataM`

Although a monadic catamorphism exists [7], this still has the same prior issues noted in Section 2.3. It lacks the ability retain type arguments of the abstract data type that is being folded.

```
cataM :: (Traversable f, Monad m) ⇒ (∀a. f a → m a) → Fix f → m a
cataM algM (In x) = algM ≪ mapM (cataM algM) x
```

To solve the problem of not retaining types, an indexed monadic catamorphism will be introduced.

However, to be able to define such a fold, a indexed monadic map is required. This will be added by extending the `IFunctor7` instance to also include a function named `imapM7`.

```
class IFunctor7 iF where
  imap7 :: (∀a b c d e f g. f' a b c d e f g → g' a b c d e f g) → iF f' a b c d e f g → iF g' a b c d e f g
  imapM7 :: Monad m ⇒ (∀a b c d e f g. f' a b c d e f g → m (g' a b c d e f g))
    → iF f' a b c d e f g
    → m (iF g' a b c d e f g)
```

Check this doesn't exist somewhere..

The definition for this new function `imapM7` for each instance closely follows the non-monadic version. Here is the definition of the new `IFunctor7` instance for `Beside`:

```
instance IFunctor7 Beside where
  imap7 f (Beside l r) = Beside (f l) (f r)
  imapM7 f (Beside l r) = do
    l' ← f l
    r' ← f r
    return (Beside l' r')
```

Now that there is a indexed monadic map, it is possible to define the a monadic catamorphism for indexed functors:

```
icataM7 :: (IFunctor7 iF, Monad m)
  => (∀ a b c d e f g. iF f' a b c d e f g → m (f' a b c d e f g))
  → IFix7 iF a b c d e f g
  → m (f' a b c d e f g)
icataM7 algM (In7 x) = algM ≪≪ imapM7 (icataM7 algM) x
```

4.4.2 BuildNetworkAlg

To use `icataM7` to fold a `Circuit` into a `BasicNetwork`, an algebra is required. However, a standard algebra will not be able to complete this transformation. Consider this example `Circuit` with two tasks executed in sequence.

```
example = task1
        <>
        task2
```

In a standard algebra both sides of the `Then` constructor would be evaluated. In this case it would produce two disjoint networks, both with their own input and output channels. The algebra for `Then`, would then need to join the output channels of `task1` with the input channels of `task2`. However, it is not possible to join channels together. Instead, the output channels from `task1` need to be accessible when creating `task2`. This is referred to as a *context-sensitive* or *accumulating* fold.

To be able to have an accumulating fold, inside an indexed catamorphism a carrier data type is required to wrap up this function. This carrier, which shall be named `N`, contains a function that when given a network that has been accumulated up to that point, then it is able to produce a network including the next layer in a circuit. The type of the layer being folded will be `Circuit a b c d e f g`.

```
newtype N n asS asT asA a b c d e f g = N
  { unN :: n asS asT asA a b c → IO (n asS asT asA d e f) }
```

This newtype has 3 additional type parameters at the beginning, namely: `asS`, `asT`, `asA`. They represent the input types to the initial circuit. Since the accumulating fold will work layer by layer from the top downwards, these types will remain constant and never change throughout the fold.

this clear enough?
maybe a diagram showing the layer and all the types

Classy Algebra An algebra type class can now be defined. This will ensure that the approach remains with modular: a new instance can be made when adding a new constructor to the language.

InitialPipes

4.4.3 The Translation

Basic Constructors Nice and simple

Task A lot nicer than you might think

Then funny step in the accumulating fold

Beside Time to unleash hell...

4.5 UUIDS

Why are they needed?

How are they added?

4.6 Failure in the Process Network

Why?

4.6.1 Maybe Monad

No error messages

4.6.2 Except Monad

based on Either

Chapter 5

Examples

5.1 Song Data Aggregation

A use case for circuits is when building data pipelines. Here there are many tasks that can only be executed when other data files have been produced. A data pipeline will also benefit from being parallel to improve run-times. A circuit can help to build a parallel data pipeline, without the user having to worry about how to combine all their data manipulations together, in a way that will not run into concurrency problems.

Consider the example, where an audio streaming service would like to create a playlist full of new songs to listen to. This could require a model that can predict your songs based on the top 10 artists and songs that you have listened to over the last 3 months. However, each of the months data is stored in different files that need aggregating together, before they can be input into the model.

5.1.1 Building the pre-processing Circuit

This circuit will need to have 3 different inputs — each months listening history. It will also have 2 outputs: the top 10 songs and artists. To begin with a dataflow diagram can be constructed — seen in Figure 5.1. This will model all the dependencies between each of the pre-processing tasks.




Figure 5.1: A dataflow diagram for pre-processing the song data

This now needs to be converted into a `Circuit`. The first step is to create all the tasks that will be used. Both the `AggSongs` and `AggArtists` tasks will require three inputs, which will be CSVs. This means that it is possible to make use of the pre-defined data store: `NamedCSVStore`. This requires the definition of `ToNamedRecord` and `FromNamedRecord` for each of the used data types, which will be omitted.

This alignment on the types.. good or nah? might help make the 3rd and 6th arg clearer?

Appendix?
or will just
being in
the code
submis-
sion be
enough?

```

aggSongsTask :: Circuit
  '[NamedCSVStore,      NamedCSVStore,      NamedCSVStore]
  '[      [Listen],      [Listen],      [Listen]]
  '[NamedCSVStore [Listen], NamedCSVStore [Listen], NamedCSVStore [Listen]]
  '[VariableStore]
  '[      [TrackCount]]
  '[VariableStore [TrackCount]]
  N3
aggSongsTask = multiInputTask f Empty
where
  f :: HList '([Listen], [Listen], [Listen]) → [TrackCount]
  f (HCons day1 (HCons day2 (HCons day3 HNil))) =
    (map (uncurry TrackCount) · reverse · count_ · map track) (day1 ++ day2 ++ day3)

```

```

aggArtistsTask :: Circuit
  '[NamedCSVStore,      NamedCSVStore,      NamedCSVStore]
  '[      [Listen],      [Listen],      [Listen]]
  '[NamedCSVStore [Listen], NamedCSVStore [Listen], NamedCSVStore [Listen]]
  '[VariableStore]
  '[      [ArtistCount]]
  '[VariableStore [ArtistCount]]
  N3
aggArtistsTask = multiInputTask f Empty
where
  f :: HList '([Listen], [Listen], [Listen]) → [ArtistCount]
  f (HCons day1 (HCons day2 (HCons day3 HNil))) =
    (map (uncurry ArtistCount) · reverse · count_ · map (artist · track)) (day1 ++ day2 ++ day3)

```

The final task to define is `Top10`, however, this task is used multiple times. Therefore, it would be beneficial if the type was polymorphic so that it can receive and input of both `[ArtistCount]` and `[TrackCount]`. This is possible to do:

```

top10Task :: (ToNamedRecord a, FromNamedRecord a, DefaultOrdered a)
  ⇒ FilePath
  → Circuit '[VariableStore] '[a] '[VariableStore [a]]
            '[NamedCSVStore] '[a] '[NamedCSVStore [a]]
            N1
top10Task filename = functionTask (take 10) (NamedCSVStore filename)

```

Now that all of the tasks have been defined, they need to be combined into a circuit. The dataflow diagram seen in Figure 5.1, will prove to be helpful. The first obvious problem is how to create a circuit that can achieve the top part of the diagram, transforming the inputs so that they can be passed into two tasks. This can be dealt with in layers. Using the diagrams associated with each constructor, it is possible to convert a layer into a composition of Circuits. The first layer, will be responsible for duplicating each of the three inputs.

```
layer1 = replicate <> replicate <> replicate
```

Scanning down the dataflow diagram the next layer to deal with is the two swaps that occur: month 1 and 2, and month 2 and 3.

```
layer2 = id <> swap <> swap <> id
```

The final layer will swap month 1 with month 3.

```
layer3 = id <> id <> swap <> id <> id
```

These layers can be stacked on top of each other to provide a full transformation:

```
transformation = layer1
                  <>
                  layer2
                  <>
                  layer3
```

The tasks can now be combined with this transformation to provide the full circuit:

```
preProcPipeline :: Circuit
  '[NamedCSVStore,      NamedCSVStore,      NamedCSVStore]
  '[      [Listen],      [Listen],      [Listen]]
  '[NamedCSVStore [Listen], NamedCSVStore [Listen], NamedCSVStore [Listen]]
  '[NamedCSVStore, NamedCSVStore]
  '[      [ArtistCount],      [TrackCount]]
  '[NamedCSVStore [ArtistCount], NamedCSVStore [TrackCount]]
  N3
preProcPipeline = transformation
                  <>
                  aggSongsTask              <> aggArtistsTask
                  <>
                  top10Task "top10Artists.csv" <> top10Task "top10Artists.csv"
```

Again it can be seen how this structure of tasks directly correlates with the dataflow diagram previously seen in Figure 5.1. This helps to make it easier when designing circuits as it can be constructed visually level by level.

5.1.2 Building the prediction Circuit

Now that there is a pipeline to pre-process the data for the model, a new playlist can be created. Again it is possible to build a circuit that combines with the pre-processing circuit.

The audio company train models for each user and store them in a cloud storage service. This would be a good use-case for creating a new `DataStore` — called a `ModelStore`. Here is a sketch for how a `ModelStore` could be defined and used.

```
data NewSongsPlaylist = NewSongsPlaylist
data ModelStore a = ModelStore a
```

A new `DataStore` instance can be defined for each different `Model`. Here is the `NewSongsPlaylist` as an example

```
instance DataStore ModelStore NewSongsPlaylist where
  -- Fetch from cloud storage and load into the program
  fetch uuid (ModelStore NewSongsPlaylist) = ...
  save  uuid (ModelStore NewSongsPlaylist) _ = error "not a valid action"
```

With this `DataStore` a new task can be defined that will, take a users top 10 songs and artists, and a model as input. It will then output a list of songs to create a playlist.

```

predictTask :: Circuit
  '[NamedCSVStore, NamedCSVStore, , ModelStore]
  '[ [ArtistCount], [TrackCount], NewSongsPlayList]
  '[NamedCSVStore [ArtistCount], NamedCSVStore [TrackCount], ModelStore NewSongsPlayList]
  '[NamedCSVStore]
  '[ [Track]]
  '[NamedCSVStore [Track]]
  N3
predictTask = multiInputTask f (NamesCSVStore "newSongsPlaylist.csv")
  where
    f :: HList '[ [ArtistCount], [TrackCount], NewSongsPlayList] → [Track]
    f (HCons topArtists (HCons topSongs (HCons model HNil))) =
      predict model topArtists topSongs

```

The `predictTask` can now be combined with the pre-processing circuit to create a circuit that is able to create a new playlist from listening history.

```

createPlaylist = preProcPipeline <> id
  <>
  predictTask

```

5.2 lhs2TeX Build System

One use case for a task based dependency system is a build system. For example, a Makefile is a way of specifying the target files from some source files, with a command that can be used to generate the target file. A Circuit could also be used to model such a system.

Consider this dissertation, which is made using \LaTeX . This project is made up of multiple sub-files, each written in a literate Haskell format. Each of these files needs to be pre-processed by the `lhs2TeX` command to produce the `.tex` source file. Once each of these files has been generated, then the \LaTeX project can be built into a PDF file.

5.2.1 Building the Circuit

A Circuit that makes use of the `mapC` operator can be used here. To do so a Circuit can be defined that is able to build a single `.tex` file from a `.lhs`. This has to make use of the standard `task` constructor:

```

buildLhsTask :: Circuit '[VariableStore] '[String] '[VariableStore String]
  '[VariableStore] '[String] '[VariableStore String]
  N1
buildLhsTask = task f Empty
  where
    f :: UUID
      → HList '[VariableStore] '[String]
      → VariableStore String
      → ExceptT SomeException IO (VariableStore String)
    f _ (HCons' (Var flnName) HNil') _ = do
      let fOutName = flnName <.> ".tex"
      lift (callCommand ("lhs2tex -o " ++ fOutName ++ " " ++ flnName))
      return (Var fOutName)

```

This Circuit makes use of the `callCommand` function from the `System.Process` library. This allows the task to execute external commands, that may not necessarily be defined in Haskell. A similar Circuit can be defined that will compile the `.tex` files and produce a PDF.

```

buildTexTask :: String → Circuit '[VariableStore] '[[String]] '[VariableStore [String]]
                                   '[VariableStore] '[String] '[VariableStore String]
                                   N1
buildTexTask name = task f Empty
where
  f :: UUID
  → HList' '[VariableStore] '[String]]
  → VariableStore String
  → ExceptT SomeException IO (VariableStore String)
  f mainFileName (HCons' (Var _) HNil') _ = do
    lift
      (callCommand
        ("texfot --no-stderr latexmk -interaction=nonstopmode -pdf "
        ++ "-no-shell-escape -bibtex -jobname="
        ++ name
        ++ " "
        ++ mainFileName
        )
      )
    return (Var "dissertation.pdf")

```

The `buildTexTask` also demonstrates how it is possible to pass a global parameter into a task. This can allow tasks to be made in a more reusable way. Saving a user from defining multiple variations of the same task.

These two tasks can now be combined into a `Circuit`. This `Circuit` can be interpreted as inputting a list of `.lhs` files, which are sequentially compiled to `.tex` files by the `mapC` operator. These files are *then* built by the `buildTexTask` to produce a PDF file as output.

```

buildDiss :: String → Circuit '[VariableStore] '[[String]] '[VariableStore [String]]
                                   '[VariableStore] '[String] '[VariableStore String]
                                   N1
buildDiss name = mapC buildLhsTask Empty
                <>
                buildTexTask name

```

5.2.2 Using the Circuit

To use this `Circuit` first a network needs to be started, then a list of `.lhs` files are input into the network. A call is then made to read the output from the network, this ensures that the network has finished compiling before the program terminates. Finally the network is stopped and all threads created are terminated.

This example makes use of a `YAML` configuration file to specify that aligns with the `Config` data type. Its purpose is to detail the files that will be included within the build.

```

data Config = Config
  { mainFile    :: FilePath
  , outputName :: String
  , lhsFiles    :: [FilePath]
  }
deriving (Generic, FromJSON, Show)
loadConfig :: IO Config
loadConfig = loadYamlSettings ["dissertation.tex-build"] [] ignoreEnv

```

```

main :: IO ()
main = do
  config ← loadConfig
  n      ← startNetwork (buildDiss (outputName config)) :: IO
    (BasicNetwork '[VariableStore] '[[String]] '[VariableStore [String]]
      '[VariableStore] '[String] '[VariableStore String])
  write (mainFile config -<. > "tex") (HCons' (Var (lhsFiles config)) HNil) n
  _ ← read n
  stopNetwork n

```

In fact, this dissertation is actually built with this system.

5.3 Types saving the day

Consider an example shown in the docs for Luigi [20], that is made up of two tasks. The first generates a list of words and saves it to a file and second counts the number of letters in each of those words. The counting letters task is dependent on the words being generated.

Figure 5.2, shows an implementation of such a system, in the Python library called Luigi. However, this implementation has a very subtle bug! **GenerateWords** writes the words to a file separated by new lines, but **CountLetters** reads that same file as a comma-separated list. This shows a key flaw in this system, it is up to the programmer to ensure that they write the outputs correctly, and then that they read that same file in the same way. This error, would not even cause a run-time error, instead, it will just produce the incorrect result. For a developer this is extremely unhelpful, it means more of time is used writing tests — something that no one enjoys.

The Fix Why not eliminate the need for all of this with **DataStores** and types. As previously mentioned in Section 3.2.1, a **DataStore** can be used to abstract the reading and writing of many different sources. This will help to ensure correctness of this step, by eliminating any possible duplicated code. Instead, just having the **fetch** and **save** methods to test.

The second greater benefit, is to use **DataStores** in combination with the type system. Each constructor for a **Circuit** will, enforce that the types of a **DataStore** align correctly. It would not be possible to feed the output of one task, with the type **FileStore [String]** into a task that expects a **CommaSepFile [String]**. The same example as before can be seen in Figure 5.3. In this example it will fail to compile, giving the error:

```
> Couldn't match type 'CommaSepFile' with 'FileStore'
```

This will benefit the user as it reduces the feedback loop of knowing if the program will succeed. Previously the whole data pipeline had to be run, whereas now this information can be informed to the user at compile-time.

```

import luigi

class GenerateWords(luigi.Task):
    def output(self):
        return luigi.LocalTarget('words.txt')

    def run(self):
        # write a dummy list of words to output file
        words = ['apple', 'banana', 'grapefruit']

        with self.output().open('w') as f:
            for word in words:
                f.write('{word}\n'.format(word=word))

class CountLetters(luigi.Task):
    def requires(self):
        return GenerateWords()

    def output(self):
        return luigi.LocalTarget('letter_counts.txt')

    def run(self):
        # read in file as list
        with self.input().open('r') as infile:
            words = infile.read().split(',')

        # write each word to output file with its corresponding letter count
        with self.output().open('w') as outfile:
            for word in words:
                outfile.write('{word}:{letter_count}\n'.format(
                    word=word,
                    letter_count=len(word)
                ))

```

Figure 5.2: A Broken Luigi Example

```

generateWords :: Circuit '[VariableStore] '[(())] '[VariableStore ()]
                    '[FileStore] '[String]] '[FileStore [String]]
                    N1
generateWords = functionTask (const ["apple", "banana", "grapefruit"]) (FileStore "fruit.txt")

countLetters :: Circuit '[CommaSepFile] '[[String]] '[CommaSepFile [String]]
                    '[FileStore] '[String]] '[FileStore [String]]
                    N1
countLetters = functionTask (map f) (FileStore "count.txt")
    where
        f word = (concat [word, ":", show (length word)])

circuit :: Circuit '[VariableStore] '[(())] '[VariableStore ()]
                    '[FileStore] '[String]] '[FileStore [String]]
                    N1
circuit = generateWords <=> countLetters

```

Figure 5.3: A Broken Circuit Example

Chapter 6

Critical Evaluation

6.1 Runtime comparison

6.1.1 Lazy evaluation problems

6.2 Use of Library

Something about how DataStores prevent the need for `luigi.ExternalTask` at the beginning.

6.2.1 Other Libraries

How does it compare?

Luigi Object-orientated approach Python library

Funflow Uses arrows to compose flows sequentially.

6.3 Type saftey

Chapter 7

Conclusion

Todo list

Change this to something meaningful	ix
Write an introduction (do near the end)	1
Add an example of simple HList	10
Add an example to get the length of a type list. Length needed later	10
Existential Types	10
Phantom type parameters	11
Monadic resource theories — Symmetric Monoidal pre-orders.	11
Motivate further why a <code>DataStore</code> needs to exist — prevents the user from reading from a source incorrectly.	13
free moniod? maybe yes it looks like what eddie did...	15
cite where this comes from?	15
maybe this is something to do with MRT??	17
Opinions on way to draw DropL/R?? or do i change it to Drop, which aligns with MRT.	17
'!': looks awful	18
Give example on the types as to why it isnt.	19
I think its a monoid in cat theory, but it isn't one in haskell. empty exists with id, but this would not conform to the types.	19
<code>DataStore</code> ' does not need <code>save</code> ' as it is never used in a task. Can only have 1 output.	19
Graphical representation of this?	21
something about the stuff Alex said	21
I can't find anything to cite that discusses this though :(, but afaik i cannot find a scenario that cannot be made...	21
??? Is there anything else that maybe I haven't achieved? memorisation?	23
flesh this out a bit more, doesnt really fit in the current context	23
cite	24
combined = co-product	25
Should I add an example of how to use it?	25
Does basic make it sound simple? should i got with standard or something like that??	25
Should I add the definition of these?	26
Check this doesn't exist somewhere..	26
this clear enough? maybe a diagram showing the layer and all the types	27
Appendix? or will just being in the code submission be enough?	29
This alignment on the types.. good or nah? might help make the 3rd and 6th arg clearer?	29

Bibliography

- [1] Quartz composer user guide, Jul 2007.
- [2] Dr Peter Bentley. The end of moore’s law: what happens next?, Apr 2020.
- [3] James Cheney and Ralf Hinze. First-class phantom types. 01 2003.
- [4] Edward Costello. *Figure 15. Connecting a Math patch to the Sprite Width and Height inlets.* Creating Graphical User Interfaces for Csound using Quartz Composer. Csound Journal, Nov 2012.
- [5] Richard A. Eisenberg, Dimitrios Vytiniotis, Simon Peyton Jones, and Stephanie Weirich. Closed type families with overlapping equations. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL ’14, page 671–683, New York, NY, USA, 2014. Association for Computing Machinery.
- [6] Richard A. Eisenberg and Stephanie Weirich. Dependently typed programming with singletons. In *Proceedings of the 2012 Haskell Symposium*, Haskell ’12, page 117–130, New York, NY, USA, 2012. Association for Computing Machinery.
- [7] M.M. Fokkinga. Monadic maps and folds for arbitrary datatypes. *Memoranda informatica*, (94-28):–, June 1994. Imported from EWI/DB PMS [db-utwente:tech:0000003538].
- [8] Felienne Hermans, Martin Pinzger, and Arie van Deursen. Breviz: Visualizing spreadsheets using dataflow diagrams, 2011.
- [9] Ralf Hinze. An algebra of scans. In Dexter Kozen, editor, *Mathematics of Program Construction*, pages 186–210, Berlin, Heidelberg, 2004. Springer Berlin Heidelberg.
- [10] Tiffany Hu. Building out the seatgeek data pipeline, Jan 2015.
- [11] Gilles Kahn. The semantics of a simple language for parallel programming. In Jack L. Rosenfeld, editor, *Information Processing, Proceedings of the 6th IFIP Congress 1974, Stockholm, Sweden, August 5-10, 1974*, pages 471–475. North-Holland, 1974.
- [12] Oleg Kiselyov, Ralf Lämmel, and Kean Schupke. Strongly typed heterogeneous collections. In *Proceedings of the 2004 ACM SIGPLAN Workshop on Haskell*, Haskell ’04, page 96–107, New York, NY, USA, 2004. Association for Computing Machinery.
- [13] S.M. Lane, S.J. Axler, Springer-Verlag (Nowy Jork)., F.W. Gehring, and P.R. Halmos. *Categories for the Working Mathematician*. Graduate Texts in Mathematics. Springer, 1998.
- [14] E. A. Lee and T. M. Parks. Dataflow process networks. *Proceedings of the IEEE*, 83(5):773–801, 1995.
- [15] Conor McBride. Functional pearl: Kleisli arrows of outrageous fortune. *Journal of Functional Programming (accepted for publication)*, 2011.
- [16] G. E. Moore. Cramming more components onto integrated circuits, reprinted from electronics, volume 38, number 8, april 19, 1965, pp.114 ff. *IEEE Solid-State Circuits Society Newsletter*, 11(3):33–35, 2006.
- [17] Simon Peyton Jones, Dimitrios Vytiniotis, Stephanie Weirich, and Geoffrey Washburn. Simple unification-based type inference for gads. *SIGPLAN Not.*, 41(9):50–61, September 2006.

- [18] Tom Schrijvers, Simon Peyton Jones, Manuel Chakravarty, and Martin Sulzmann. Type checking with open type functions. In *Proceedings of the 13th ACM SIGPLAN International Conference on Functional Programming*, ICFP '08, page 51–62, New York, NY, USA, 2008. Association for Computing Machinery.
- [19] Spotify. Luigi: <https://github.com/spotify/luigi>.
- [20] Spotify. Tasks, Apr 2020.
- [21] Josef Svenningsson and Emil Axelsson. Combining deep and shallow embedding of domain-specific languages. *Computer Languages, Systems & Structures*, 44:143–165, 2015. SI: TFP 2011/12.
- [22] WOUTER SWIERSTRA. Data types à la carte. *Journal of Functional Programming*, 18(4):423–436, 2008.
- [23] Tableau. Tableau prep builder & prep conductor: A self-service data preparation solution, 2021.
- [24] Phillip Wadler. The expression problem, Nov 1998.
- [25] Jamie Willis, Nicolas Wu, and Matthew Pickering. Staged selective parser combinators. *Proc. ACM Program. Lang.*, 4(ICFP), August 2020.
- [26] Nicolas Wu. Yoda: A simple combinator library, 2018.
- [27] Brent A. Yorgey, Stephanie Weirich, Julien Cretin, Simon Peyton Jones, Dimitrios Vytiniotis, and José Pedro Magalhães. Giving haskell a promotion. In *Proceedings of the 8th ACM SIGPLAN Workshop on Types in Language Design and Implementation*, TLDI '12, page 53–66, New York, NY, USA, 2012. Association for Computing Machinery.